Harshith Kadaru

Arlington, VA | kadaruharshith@gmail.com | 571-571-3336 | kadaru242003.github.io | GitHub | LinkedIn

EDUCATION

The George Washington University — Washington, DC
Master of Science, Data Science, GPA: 4.0 — Expected May 2027

CVR College of Engineering — Hyderabad, India
Bachelor of Technology, Computer Science and Information Technology — April 2025

SKILLS

Languages: Python, SQL, R, TypeScript, JavaScript
ML and Statistics: scikit-learn, LightGBM, PyTorch, pandas, SciPy, time series forecasting, A/B testing, power analysis
AI and LLMs: LLM evaluation, LLM-as-a-judge, agent harnesses, MCP servers, tool calling, RAG, embeddings, ChromaDB
Engineering: FastAPI, React, Redis, MySQL, MongoDB, Docker, GCP (Cloud Run, Cloud SQL), GitHub Actions CI

EXPERIENCE

MATClinics (SUDx Analytics) — Maryland (Hybrid)
Software Engineering Intern — Jun 2026 - Aug 2026

- Migrated two production ML scoring pipelines from R / H2O / Java to Python, with 100% and 97.3% parity.
- Root-caused 5 months of silent job failures, backfilled the missing scores, and made the monthly rebuild idempotent.
- Moved real-time updates from in-memory WebSockets to Redis pub/sub to work across scaled Cloud Run instances.
- Built Python ETL into MySQL (Cloud SQL) and FastAPI telehealth services (Google Meet, Twilio SMS) with audit logging.
- Shipped 15+ React / TypeScript features to a multi-site clinical dashboard from direct provider feedback.

The George Washington University — Washington, DC
Graduate Instructional Assistant, DATS 1001: Data Science for All — Aug 2026 - Present

- Lead labs, grading, and office hours for R-based data wrangling, visualization, and statistical inference.

PROJECTS

claimcheck | Python, Docker, GitHub Actions | GitHub — In Progress
Open-source benchmark testing whether AI coding agents falsely claim success

- Built 205 tasks with verified ground truth and a harness running GPT-OSS 120B and 20B through 530 runs at $0.
- Hardened grading against 8 ways to fake a pass; 12 cheat strategies (1,236 attempts) get none through.
- Built an LLM-as-a-judge from a different model family and pre-registered the analysis (McNemar and Fisher tests).
- Caught GPT-OSS 120B reverse-engineering test data to pass an impossible task, then quietly claiming success.

ML Model Service with Drift Monitoring | Python, LightGBM, FastAPI, Docker | GitHub — Oct 2026

- Served a LightGBM hotel cancellation model (ROC AUC 0.878) with FastAPI in Docker, logging predictions to SQLite.
- Replayed 55,095 later bookings: AUC fell to 0.76, but drift alarms did no better than random (precision 0.29).
- Traced the drop to a shift in booking mix (Spearman -0.84) and measured detector power and false alarms by simulation.

claimcheck-mcp | Python, MCP, Docker | GitHub — Oct 2026

- MCP server giving AI agents 3 tools over the 205 tasks in a Docker sandbox; fixed a symlink leak of answer files.

A/B Experimentation and Decision Framework | Python, SciPy, pandas | GitHub — Personal

- Built SRM checks, z-tests, Welch's t-tests, and guardrails that decide ship, continue, or rollback for each experiment.
- Showed with 2,000 A/A simulations that peeking raised false positives from 5% to 19%; O'Brien-Fleming restored 5%.
- Added power analysis (sample size, achieved power, test duration) validated against statsmodels, with 92 tests in CI.

Demand Forecasting with Backtesting | Python, LightGBM, statsmodels | GitHub — Oct 2026

- Forecast weekly sales for 3,049 Walmart (M5) items with ETS and LightGBM using lag and calendar features.
- 6-fold rolling-origin backtests with leakage tests: LightGBM MASE 0.963 vs seasonal naive 1.437.

RiskBot | Python, FastAPI, ChromaDB, GPT-OSS 120B | GitHub | Live Demo — Dec 2025

- Flags sensitive HR and personal data across whole files, with document Q&A using top-k retrieval in ChromaDB.

ACHIEVEMENTS

- **Winner, GW Data Science Association Datathon 2026:** built a 24-page Python and Streamlit platform analyzing 500K+ rows, using K-Means clustering, regression, and simulation to recommend policy interventions.